

Testing Battle.net

(Before deploying to millions of players)

Ben Deane

Principal Engineer, Blizzard Entertainment

bdeane@blizzard.com, @ben_deane

13th May 2015

Battle.net infrastructure

- About 750,000 lines of C++
 - Servers + client libraries
- "Battle.net Game Service"
 - Authenticate players
 - Social: friends, presence
 - Matchmaking (cooperative/competitive)
 - Achievements/profiles

Battle.net is highly...

- Distributed
- Asynchronous
- Configured
- Fault-prone
- Architecture-varied
 - inheritance
 - composition
 - value-oriented

Battle.net integration testing

gameservice

Battle.net Game Service

Updated 3 minutes ago

C# ★ 52 📄 89

- API testing is pretty robust
 - and certainly valuable
- But this doesn't help me in the moment
 - slow to build
 - slow to run
 - needs a full environment

My journey towards effective unit testing

- No practice at unit testing
- Large project with many moving parts
- Mature lower level libraries
- New code (features) added at an alarming rate

What's typically well-tested?

- UTF-8 string conversion
- String interpolation
- URL parsing/decomposition
- Stats/math code

These things are "easy mode" for tests.

Not-so-well tested?

- Filesystem interaction (caching downloaded objects)
- Matchmaking algorithms
- Queueing/Load balancing algorithms

These things are harder to test. Where to start?

My conclusions

- We don't do unit testing because we aren't practised at it
- Because we don't do "TDD", we *can't* do unit testing
 - Legacy code is poorly structured
- We have a test framework

My goals for "unit" tests

- Fast
- No data/process dependencies
- Automated
- Binary pass/fail
- Independent
- No test-only interface support
- By me, for me

No magic bullet

- I wrote a load of mocks
- Set up a lot of data structures for test
- A lot of testing code to keep bug-free
- But along the way I found
 - better code structure
 - useful techniques

Enemies of testing

- Global state
- Deep inheritance
- Mixing concerns, coupling
- I/O

Enemies of testing

- Doing work in constructors (cf RAI)
- Lack of dependency injection
- Wide interfaces (especially when passed to constructors)

Exhibit A: hard to test

```
class ChannelBase : public rpc::Implementor<protocol::channel::Channel>;
class ChannelImpl : public ChannelBase;

class PresenceChannelImpl : public ChannelImpl
{
public:
    PresenceChannelImpl(
        Process* process,
        rpc::RPCDispatcher* insideDispatcher,
        const EntityId& entityId,
        ChannelDelegate* channelDelegate,
        ChannelOwner* owner,
        const PresenceFieldConfigMap& fieldMap);
};
```

Exhibit A: hard to test

```
class ChannelBase : public rpc::Implementor<protocol::channel::Channel>;
class ChannelImpl : public ChannelBase;

class PresenceChannelImpl : public ChannelImpl
{
public:
    PresenceChannelImpl(
        Process* process,
        rpc::RPCDispatcher* insideDispatcher,
        const EntityId& entityId,
        ChannelDelegate* channelDelegate,
        ChannelOwner* owner,
        const PresenceFieldConfigMap& fieldMap);
};
```

Exhibit A: hard to test

```
class ChannelBase : public rpc::Implementor<protocol::channel::Channel>;
class ChannelImpl : public ChannelBase;

class PresenceChannelImpl : public ChannelImpl
{
public:
    PresenceChannelImpl(
        Process* process,
        rpc::RPCDispatcher* insideDispatcher,
        const EntityId& entityId,
        ChannelDelegate* channelDelegate,
        ChannelOwner* owner,
        const PresenceFieldConfigMap& fieldMap);
};
```

Exhibit B: hard to test

```
class AchievementsServiceImpl
: public bnet::achievements::AchievementsService
, public AchievementsServiceStaticDataLoader
{
public:
    AchievementsServiceImpl(
        bnet::internal::ServerHelper& serverHelper,
        mysql::Databases* mysql);
};
```


Exhibit B: hard to test

```
class AchievementsServiceImpl
    : public bnet::achievements::AchievementsService
    , public AchievementsServiceStaticDataLoader
{
public:
    AchievementsServiceImpl(
        bnet::internal::ServerHelper& serverHelper,
        mysql::Databases* mysql);
};
```

Exhibit B: hard to test


```
class ServerHelper
{
public:
    ServerHelper(...); // 12 args!

    rpc::RPCServer* GetInsideRPCServer() const;
    rpc::RPCServer* GetOutsideRPCServer() const;
    ...
};
```

In hindsight, this was a mistake


Class structure for testing

- Base class (contains logic)
- Derived class (contains I/O, config, etc)




Class structure for testing

- Component class (contains logic)
- Entity/Object class (contains I/O, config, etc)



Example: Queueing for games



Queueing for games

GameQueueBase contains the queueing logic

```
class GameQueueBase
{
public:
    GameQueueBase(
        shared_ptr<ServerPoolInterface> interface,
        const PopCallback& popCb,
        const UpdateCallback& updateCb,
        const PollTimerCallback& pollTimerCb,
        const NotificationTimerCallback& notificationTimerCb);

    bool    Push(...);
    size_t  Pop(...);
    void    Remove(...);
    size_t  PollQueue(...);

    ...
};
```

Queueing for games

GameQueueBase contains the queueing logic

```
class GameQueueBase
{
public:
    GameQueueBase(
        shared_ptr<ServerPoolInterface> interface,
        const PopCallback& popCb,
        const UpdateCallback& updateCb,
        const PollTimerCallback& pollTimerCb,
        const NotificationTimerCallback& notificationTimerCb);

    bool    Push(...);
    size_t  Pop(...);
    void    Remove(...);
    size_t  PollQueue(...);

    ...
};
```

Queueing for games

GameQueueBase contains the queueing logic

```
class GameQueueBase
{
public:
    GameQueueBase(
        shared_ptr<ServerPoolInterface> interface,
        const PopCallback& popCb,
        const UpdateCallback& updateCb,
        const PollTimerCallback& pollTimerCb,
        const NotificationTimerCallback& notificationTimerCb);

    bool    Push(...);
    size_t  Pop(...);
    void    Remove(...);
    size_t  PollQueue(...);

    ...
};
```


Queueing for games

GameQueueImpl deals with protocols

```
class GameQueueImpl
    : public GameQueueBase
    , public protocol::game_queue::GameQueue
{
public:
    // protocol handler functions
    virtual void AddToQueue(...);
    virtual void RemoveFromQueue(...);
    ...

    // system events
    bool OnInit(...);
    bool OnFlush(...);
    void OnShutdown(...);
    void OnPeerDisconnected(...);
    ...
};
```

Queueing for games

GameQueueImpl deals with protocols

```
class GameQueueImpl
    : public GameQueueBase
    , public protocol::game_queue::GameQueue
{
public:
    // protocol handler functions
    virtual void AddToQueue(...);
    virtual void RemoveFromQueue(...);
    ...

    // system events
    bool OnInit(...);
    bool OnFlush(...);
    void OnShutdown(...);
    void OnPeerDisconnected(...);
    ...
};
```

Queueing for games

GameQueueImpl deals with system events

```
class GameQueueImpl
    : public GameQueueBase
    , public protocol::game_queue::GameQueue
{
public:
    // protocol handler functions
    virtual void AddToQueue(...);
    virtual void RemoveFromQueue(...);
    ...

    // system events
    bool OnInit(...);
    bool OnFlush(...);
    void OnShutdown(...);
    void OnPeerDisconnected(...);
    ...
};
```

Queueing for games

GameQueueImpl deals with config

```
class GameQueueImpl
    : public GameQueueBase
    , public protocol::game_queue::GameQueue
{
public:
    ...

    // setup/config
    bool ProcessProgramConfig(...);

    // queue polling
    void StartPollTimer(...);
    void ServicePollTimer(...);
    void StartNotificationPollTimer(...);
    void ServiceNotificationPollTimer(...);
    ...
};
```

Queueing for games


GameQueueImpl deals with polling logic

```
class GameQueueImpl
    : public GameQueueBase
    , public protocol::game_queue::GameQueue
{
public:
    ...

    // setup/config
    bool ProcessProgramConfig(...);

    // queue polling
    void StartPollTimer(...);
    void ServicePollTimer(...);
    void StartNotificationPollTimer(...);
    void ServiceNotificationPollTimer(...);
    ...
};
```

Example: Matchmaking



Matchmaking

GameFactory contains matchmaking logic

```
class GameFactory
{
public:
    GameFactory(const AttributeValue& version,
                const ProgramId& programId,
                GameFactoryId id);

    virtual bool Configure(const GameFactoryConfig& config);

    ...
    virtual Error RegisterPlayers(...);
    virtual bool UnregisterPlayers(...);
    virtual Error JoinGame(...);
    ...
};
```

Matchmaking

GameFactory contains matchmaking logic

```
class GameFactory
{
public:
    GameFactory(const AttributeValue& version,
                const ProgramId& programId,
                GameFactoryId id);

    virtual bool Configure(const GameFactoryConfig& config);

    ...
    virtual Error RegisterPlayers(...);
    virtual bool UnregisterPlayers(...);
    virtual Error JoinGame(...);
    ...
};
```


Matchmaking

GameFactory contains matchmaking logic

```
class GameFactory
{
public:
    GameFactory(const AttributeValue& version,
                const ProgramId& programId,
                GameFactoryId id);

    virtual bool Configure(const GameFactoryConfig& config);

    ...
    virtual Error RegisterPlayers(...);
    virtual bool UnregisterPlayers(...);
    virtual Error JoinGame(...);
    ...
};
```

Matchmaking

GameMasterImpl deals with interactions

```
class GameMasterImpl
{
public:
    ...
    void OnPeerDisconnected(...);
    ...
    void InstantiateFactories(...);
    ...
    virtual void ListFactories(...);
    virtual void JoinGame(...);
    virtual void FindGame(...);
    virtual void GameEnded(...);
    virtual void PlayerLeft(...);
    ...
};
```

Matchmaking

GameMasterImpl deals with interactions

```
class GameMasterImpl
{
public:
    ...
    void OnPeerDisconnected(...);
    ...
    void InstantiateFactories(...);
    ...
    virtual void ListFactories(...);
    virtual void JoinGame(...);
    virtual void FindGame(...);
    virtual void GameEnded(...);
    virtual void PlayerLeft(...);
    ...
};
```

Matchmaking

GameMasterImpl deals with interactions

```
class GameMasterImpl
{
public:
    ...
    void OnPeerDisconnected(...);
    ...
    void InstantiateFactories(...);
    ...
    virtual void ListFactories(...);
    virtual void JoinGame(...);
    virtual void FindGame(...);
    virtual void GameEnded(...);
    virtual void PlayerLeft(...);
    ...
};
```

A successful pattern

- Decouple logic from other concerns
 - Dependency injection for config etc
 - Makes the logic testable
- This can be fairly easily applied even to monolithic classes
 - Just apply the inheritance pattern
 - Some testing beats no testing

Goals for testable classes

Dependency injection is probably the biggest factor affecting whether or not code *is testable at all*.

Even with DI, classes are *onerous to test* unless constructors take few arguments, using narrow interfaces.

Testing Performance/Efficiency

- Different solutions for
 - thousands (performance)
 - millions (performance + algorithms)
 - billions (algorithms by construction)
- Battle.net's working sets are in the millions

Problems in million-land

- Computations can run on a single machine
- Data structures are important to performance
 - Caching concerns, optimizations can get you 100x
 - But they can't get you 100,000x
- Algorithms are important to efficiency

Testing for performance

- Timed tests are easy, not so useful
- My machine is a Windows desktop
- Production machine is a CentOS blade
- Timed tests
 - compare times when optimizing
 - can't tell me if code is fast enough in an absolute sense

Efficiency: easy to lose

- Team of engineers hacking away on features
- $O(\log n)$ or less is required
- Easy to accidentally turn it into $O(n)$ (or worse)
- I need a way to test for algorithmic efficiency

Testing for efficiency

- A simple idea
- Run the same test with different sized inputs
- Compute ratio of times

$T_1 = (\text{time for run on data of size } N)$

$T_2 = (\text{time for run on data of size } kN)$

Bucketing

$$O(1) \Rightarrow \frac{T_2}{T_1} = 1$$

$$O(\log n) \Rightarrow \frac{T_2}{T_1} = 1 + \frac{\log(k)}{\log(N)}$$

$$O(n) \Rightarrow \frac{T_2}{T_1} = k$$

$$O(n \log n) \Rightarrow \frac{T_2}{T_1} = k \left(1 + \frac{\log(k)}{\log(N)} \right)$$

$$O(n^2) \Rightarrow \frac{T_2}{T_1} = k^2$$

This sounds easy, but. . .

- Timing is hard
 - sensitive to machine load
 - sensitive to caching effects (CPU/OS)
 - sensitive to timing function: granularity/perf
- Statistical mitigation
- Somewhat careful choice of k , N
 - I settled on 32 for each ($N = 32, kN = 1024$)

OK, but...

Where do you get different-sized inputs?
You can let the test make them...

```
const int MULT = 32;
const int N = 32;
...
// run 1 - with size N
auto sampleTime1 = test->Run(N);
test->Teardown();

test->Setup();
// run 2 - with size kN
auto sampleTime2 = test->Run(N * MULT);
...
```

OK, but...

Where do you get different-sized inputs?
You can let the test make them...

```
const int MULT = 32;
const int N = 32;
...
// run 1 - with size N
auto sampleTime1 = test->Run(N);
test->Teardown();

test->Setup();
// run 2 - with size kN
auto sampleTime2 = test->Run(N * MULT);
...
```

Let the test make them?

- Affects the timing if done naively (i.e. wrongly)
 - Adds an $O(n)$ component to the test
 - So move the timing code inside the test also
- Boilerplate in test code
- It's not ideal...

Let the test make them?

Result: a typical test

- ~20 lines setup
- ~20 lines size-related setup
- ~10 lines timing
- ~5 lines actual logic
- ~5 lines test macros

Yuck.

Let the test make them?

- It works well enough to give me confidence
 - Matchmaking won't blow up with a million players
- So I lived with this for a while...
- But I'm lazy, I don't want to maintain all this code
- And I'm a student of Haskell...

Wish-driven development

What I have

```
DEF_TEST(TestName, Suite)
{
    ...
    return test_result;
}
```

What I want

```
DEF_PROPERTY(TestName, Suite, const string& s)
{
    // do something with s
    // that should be true for any input
    ...
    return property_holds;
}
```

Test macros expand into functions

Macro...

```
DEF_PROPERTY(TestName, Suite, const string& s)
{
    ...
}
```

Expands to...

```
struct NonceStruct
{
    ...
    bool operator()(const string&);
};
bool NonceStruct::operator()(const string& s)
{
    ...
}
```

Discover the type of the function argument

Simple function_traits template

```
template <typename T>
struct function_traits
    : public function_traits<decltype(&T::operator()))>
{};
```

```
template <typename R, typename A>
struct function_traits<R(A)>
{
    using argType = A;
};
```

```
template <typename C, typename R, typename A>
struct function_traits<R(C::*)(A)>
    : public function_traits<R(A)>
{};
```

```
...
```

Discover the type of the function argument

Simple function_traits template

```
template <typename T>
struct function_traits
    : public function_traits<decltype(&T::operator()))>
{};
```

```
template <typename R, typename A>
struct function_traits<R(A)>
{
    using argType = A;
};
```

```
template <typename C, typename R, typename A>
struct function_traits<R(C::*)(A)>
    : public function_traits<R(A)>
{};
```

...

Discover the type of the function argument

Simple function_traits template

```
template <typename T>
struct function_traits
    : public function_traits<decltype(&T::operator())>
{};
```

```
template <typename R, typename A>
struct function_traits<R(A)>
{
    using argType = A;
};
```

```
template <typename C, typename R, typename A>
struct function_traits<R(C::*)(A)>
    : public function_traits<R(A)>
{};
```

```
...
```

Implement a Run function

Run() for a property test

```
// DEF_PROPERTY(TestName, Suite, TYPE) becomes...
struct NonceStruct : public Test
{
    ...
    virtual bool Run() override
    {
        // Property will type-erase NonceStruct, discover its argument type
        Property p(*this);
        // check() generates arguments to call NonceStruct(TYPE)
        return p.check();
    }
    ...
};
```


Property type-erases NonceStruct

Inside Property

```
template <typename T>
struct Internal : public InternalBase
{
    ...

    using paramType = std::decay_t<typename function_traits<T>::argType>;

    virtual bool check()
    {
        ...
        // generate a value of the right type
        paramType p = Arbitrary<paramType>::generate(...);
        // feed it to the struct's operator()
        return m_t(p);
    }

    T m_t;
};
```

Property type-erases NonceStruct

Inside Property

```
template <typename T>
struct Internal : public InternalBase
{
    ...

    using paramType = std::decay_t<typename function_traits<T>::argType>;

    virtual bool check()
    {
        ...
        // generate a value of the right type
        paramType p = Arbitrary<paramType>::generate(...);
        // feed it to the struct's operator()
        return m_t(p);
    }

    T m_t;
};
```

How to generate TYPE?

Use a template, naturally

```
template <typename T>
struct Arbitrary
{
    static T generate(size_t /*generation*/, unsigned long int /*seed*/)
    {
        return T();
    }
};
```

And specialize...

Specializing Arbitrary<T>

- Easy to write Arbitrary<T> for fundamental types
- Front-load likely edge cases
 - 0
 - `numeric_limits<T>::min()`
 - `numeric_limits<T>::max()`
- Otherwise use uniform distribution over range

Specializing Arbitrary<T>

For int-like types

```
static int generate(size_t g, unsigned long int seed)
{
    switch (g)
    {
        case 0: return 0;
        case 1: return std::numeric_limits<T>::min();
        case 2: return std::numeric_limits<T>::max();
        default:
        {
            std::mt19937 gen(seed);
            std::uniform_int_distribution<T> dis(
                std::numeric_limits<T>::min(), std::numeric_limits<T>::max());
            return dis(gen);
        }
    }
}
```

Specializing Arbitrary<T>

For int-like types

```
static int generate(size_t g, unsigned long int seed)
{
    switch (g)
    {
        case 0: return 0;
        case 1: return std::numeric_limits<T>::min();
        case 2: return std::numeric_limits<T>::max();
        default:
        {
            std::mt19937 gen(seed);
            std::uniform_int_distribution<T> dis(
                std::numeric_limits<T>::min(), std::numeric_limits<T>::max());
            return dis(gen);
        }
    }
}
```

Specializing Arbitrary<T>

- Once we have Arbitrary<T> for fundamental types...
- Easy to write for compound types
 - vector<T> etc
 - generate works in terms of generate on the contained type
 - the power of ADTs!

Specializing Arbitrary<T>

For compound types (eg vector)

```
static vector<T> generate(size_t g, unsigned long int seed)
{
    vector<T> v;
    size_t n = 10 * ((g / 100) + 1);
    v.reserve(n);
    std::generate_n(
        std::back_inserter(v), n, [&] () {
            return Arbitrary<T>::generate(g++, seed++); });
    return v;
}
```


Now we have property tests

- Macro expands `NonceStruct` with `operator()`
- Property type-erases `NonceStruct`
- `Property::Check` does:
 - `function_traits` discovery of the argument type `T`
 - `Arbitrary<T>::generate` to make a `T`
 - Call `NonceStruct::operator()`
- And plumb through parameters like number of checks, random seed

Better checks for compound types

When a check fails, find a minimal failure case

```
template <typename T>
struct Arbitrary
{
    static std::vector<T> shrink(const T& /*t*/)
    {
        return std::vector<T>();
    }
};
```

shrink returns a vector of "reduced" T's

Better checks for compound types

A simple binary search

```
static std::vector<std::basic_string<T>> shrink(  
    const std::basic_string<T>& t)  
{  
    std::vector<std::basic_string<T>> v;  
    if (t.size() < 2)  
        return v;  
    auto l = t.size() / 2;  
    v.push_back(t.substr(0, l));  
    v.push_back(t.substr(l));  
    return v;  
}
```

Call `shrink` repeatedly to find a minimal fail case

A short demo

(Demo)

Testing for efficiency (again)

Now the computer can generate N , kN values

```
static vector<T> generate(size_t g, unsigned long int seed)
{
    vector<T> v;
    size_t n = 10 * ((g / 100) + 1);
    v.reserve(n);
    std::generate_n(
        std::back_inserter(v), n, [&] () {
            return Arbitrary<T>::generate(g++, seed++); });
    return v;
}
```

Add generate_n as a tighter form of generate

Testing for efficiency (again)

Now the computer can generate N , kN values

```
static vector<T> generate_n(size_t g, unsigned long int seed)
{
    vector<T> v;
    // use g directly instead of a "loose" value
    v.reserve(g);
    std::generate_n(
        std::back_inserter(v), g, [&] () {
            return Arbitrary<T>::generate_n(g, seed++); });
    return v;
}
```

Add generate_n as a tighter form of generate

Now I can write

A sample complexity test

```
DEF_COMPLEXITY_PROPERTY(TestName, Suite, ORDER_N, const string& s)
{
    // something that's supposed to be order N...
    std::max_element(s.begin(), s.end());
}
```

And specialize Arbitrary for my own types as necessary
Much less boilerplate to maintain

So that's where I am now

- Dependency injection (little work in constructors)
- Separate logic from interaction (even in monolithic classes)
- Regular tests for "normal, identified" cases
- Timed tests when I'm optimizing
- Property-based tests for invariants
- Algorithmic complexity tests for scalability confidence

The future?

- Arbitrary opens the door for fuzz testing?
- Alternative walk strategies through the input space
 - Hilbert?
 - Morton
 - etc
- I'm still lazy; the computer isn't doing enough for me yet

Battle.net is still highly...

- Distributed
- Asynchronous
- Configured
- Fault-prone
- Architecture-varied

But more parts of it are well-tested before they leave a developer's machine.

And I'm more confident changing code with a guarantee that correctness/efficiency/scalability won't be affected.

Thanks for listening

Code: `https://github.com/elbeno/testinator`

Me: `bdeane@blizzard.com`, `@ben_deane`